48-Hour Intelligence Brief

AI Agents & Agentic AI — Multi-Agent Frameworks Landscape

October 2026 | $497 Fixed-Scope Deliverable | Investigative Cell

1. Direct Answer

The multi-agent framework landscape in late 2026 is dominated by LangGraph for production-grade, stateful, auditable systems and CrewAI for rapid role-based prototyping and business workflows. AutoGen has been placed in maintenance mode by Microsoft; its successor is the Microsoft Agent Framework. OpenAI Swarm is effectively discontinued in favor of the OpenAI Agents SDK. Production adoption estimates place LangGraph at roughly 38% of multi-agent deployments, with custom orchestration still significant and CrewAI trailing in production maturity despite strong demo velocity.

2. Key Evidence

LangGraph: Highest production footprint. Explicit stateful graphs, first-class checkpointing, human-in-the-loop interrupts, and mature LangSmith observability. Approximately 71 million monthly PyPI installs. Best for long-running, auditable, branching workflows. Pricing: free library + LangSmith Plus at $39/seat/month.

CrewAI: Fastest for role-based crews (researcher, writer, reviewer). Claims 450 million+ agentic workflows per month and strong enterprise features (FedRAMP, VPC, SSO). Approximately 20 million monthly installs. Ideal for prototypes and business process automation. Free core + limited free hosted executions; enterprise custom.

AutoGen: Maintenance mode since approximately October 2025. No major Python releases in hundreds of days. Microsoft directs users to Microsoft Agent Framework (GA April 2026). Still referenced in research but not recommended for new production builds.

Others: Microsoft Agent Framework (successor path), OpenAI Agents SDK, Google ADK, n8n, and various no-code platforms. Custom orchestration remains approximately 28% of production deployments.

3. Strongest Counterpoints

CrewAI's "Fortune 500 / 450M workflows" figures are self-reported; independent production adoption data places it well behind LangGraph. LangGraph's production advantage comes with a steeper learning curve and potential operational costs (checkpoint storage growth in Postgres is a documented issue). Many "AI agent platforms" in marketing materials are thin wrappers or no-code tools rather than true multi-agent orchestration frameworks. Pricing units vary widely. Market estimates for production share are directional, not audited census data.

4. Residual Uncertainties

Exact enterprise revenue and true production deployment counts for each framework remain vendor-influenced. Rapid consolidation continues. The landscape can shift again within 3–6 months. Real-world failure modes and long-term maintenance costs are under-documented outside of selective case studies.

5. Actionable Recommendations

1. For production systems needing control, auditability, and human approval → Default to LangGraph.

2. For rapid internal prototypes or role-based business workflows → Start with CrewAI, then migrate critical paths if needed.

3. Avoid new builds on AutoGen; evaluate Microsoft Agent Framework only if already deep in the Microsoft stack.

4. Always calculate total cost of ownership (framework + observability + model API + storage + engineering time).

5. Validate with a narrow pilot on your actual workflow before committing.

This is a sample deliverable matching the $497 48-Hour Intelligence Brief format. Concise, source-linked, challenge-tested, and decision-oriented. Prepared by the Investigative Cell.